

Differential and Rotational Cryptanalysis of Round-reduced MORUS

Ashutosh Dhar Dwivedi¹, Paweł Morawiecki¹ and Sebastian Wójtowicz¹

¹*Institute of Computer Science, Polish Academy of Sciences, Warsaw, Poland*
{ashudhar7, pawel.morawiecki, sebastian.wojtowicz}@gmail.com

Keywords: authenticated encryption, MORUS, rotational cryptanalysis, internal differential cryptanalysis, CAESAR competition

Abstract: In this paper we investigate the security margin of MORUS — an authenticated cipher taking part in the CAESAR competition. We propose a new key recovery approach, which can be seen as an accelerated exhaustive search. We also verify the resistance of MORUS against internal differential and rotational cryptanalysis. Our analysis reveals that the cipher has a solid security margin and a lack of round constants does not bring any weakness. Our work helps to reliably evaluate this new, high-performance algorithm, which is particularly important in the context of the ongoing CAESAR competition.

1 INTRODUCTION

A cryptographic algorithm which provides both confidentiality and authenticity, is called an authenticated encryption (AE) or simply an authenticated cipher. It encrypts and authenticates messages using both a secret key (shared by the sender and the receiver) as well as a public number (called a nonce). AE algorithms are often built as various combinations of block ciphers, stream ciphers, message-authentication codes and hash functions.

The great interest and importance of AE have been manifested by the announcement of a new public call for AE algorithms — the CAESAR competition (CAE,). The contest has started in 2014 and has received worldwide attention. In the first round, 57 algorithms were submitted and now (January 2017), in the third round, 16 ciphers are still in the race. Particularly MORUS (Wu and Huang,) — the cipher we focus on — has advanced to the third round of the CAESAR competition. MORUS exhibits excellent software performance and can be also a great choice for hardware platforms. However, the problem is very little third-party cryptanalysis of this algorithm (see Related work below) as the third-party investigation is essential to build trust towards new, promising ciphers.

MORUS has an interesting design feature, namely a lack of round constants. Typically, the round constants are introduced to break a symmetry between rounds and/or between parts of the state. The MORUS designers, however, chose a different ap-

proach. A pseudorandom constant is introduced only into the initial state and with a strong initialization phase all possible state symmetries should be eliminated. In our work, we want to verify whether such design decisions lead to a weakness exploitable by internal differentials or rotational cryptanalysis. These two techniques were previously used to attack some algorithms (or their reduced variants), where the state symmetries were not fully disturbed by the round constants (Dinur et al., 2013; Morawiecki et al., 2013).

Related work As mentioned, there are very few cryptanalytic works on MORUS. Mileva *et al.* made some observations and they described the distinguisher (in a nonce-reuse scenario) (Mileva et al., 2015). However, the relevance of this work to the cipher’s security was questioned by the MORUS designers (MORUSGoogleGroupDiscussin,). Some concerns about MORUS security is highlighted in (Saarinen, 2016), where the author claims that MORUS represents significantly elevated adaptive-chosen-plaintext attack risk. Full round SAT analysis was made by Dwivedi *et al.* on MORUS-640 (Dwivedi et al., 2016).

In our work, besides the classic differential cryptanalysis, we apply internal differentials and rotational cryptanalysis. Internal differentials were first proposed by Peyrin (Peyrin, 2010) in the attack on the Grøstl hash function. In (Dinur et al., 2013), Dinur et al. showed that internal differentials could be also used to produce collisions against the round-reduced Keccak, exploiting the round constants and their re-

lation to the state symmetry. Most recently, internal differentials were combined with the boomerang technique, which led to the 8-round practical distinguisher of the Keccak-f permutation (Jean and Nikolic, 2015).

Rotational cryptanalysis was formally introduced in (Khovratovich and Nikolić, 2010) but the technique itself was mentioned and applied earlier (Bernstein, 2005; Knudsen et al., 2009; Standaert et al., 2006). Rotational cryptanalysis was also combined with the rebound attack and applied to the compression function of the SHA-3 candidate Skein and its underlying cipher Threefish (Khovratovich et al., 2010). In (Morawiecki et al., 2013), the preimage attack against the round-reduced Keccak (Bertoni et al., b) is given and the attack takes advantage of the low Hamming weight round constants.

Our contribution The classic techniques for the key recovery (such as so called $R - 1$ differential attack) cannot be directly applied to MORUS due to a different structure of the algorithm and limitations introduced by the nonce requirement. We propose a *new key recovery approach*, where differential characteristics are used in a different manner than in typical attacks known for block ciphers. Our approach can be viewed as the *accelerated exhaustive search* and can be combined with other techniques such as rotational cryptanalysis.

To quantify the security margin and gain more insight into the ciphers, we analyse reduced variants. Specifically, we reduce a number of rounds (steps) in the initialization phase, all the other parameters of the algorithms are the same as in the full variants. Our best result is the theoretical key recovery attack on MORUS-1280-256, where the initialization phase is reduced to 18 rounds. We also apply internal differential and rotational cryptanalysis to investigate whether a lack of round constants could pose a threat to the security of MORUS. The presented attacks respect the nonce requirement and need very few plaintext-ciphertext pairs. Table 1 shows our main findings.

With 80 rounds in the MORUS initialization, the algorithm has a very solid security margin against rotational and internal differential cryptanalysis. Our analysis supports the designers' claims on a strong security level of the ciphers and offers some cryptanalytic insight, which is very important for reliable evaluation of this new promising design. We argue that our findings, such as 18-round differential characteristics with probability 1, could serve as a starting point for other analysis, for example the differential-linear attack.

2 MORUS DESCRIPTION

MORUS has been designed to achieve a great performance on both software and hardware, across various platforms. Actually, MORUS is a family of three authenticated ciphers: MORUS-640-128, MORUS-1280-128, and MORUS-1280-256.

The first number in the cipher's name is a size of the state (in bits) and the second number stands for the key size. The state consists of 5 registers (either 128- or 256-bit) and each register is divided into 4 words.

The state update function consists of 5 very similar rounds and it relies on four simple bitwise operations. These operations are: AND, XOR, and two kinds of rotations. The symbol \lll denotes the rotation in the register by a given number of bits, whereas $Rotl(s, b)$ rotates words in the register s over b bits. Figure 1 shows the state update function. In our analysis we focus on MORUS with a 1280-bit state, then let us describe the initialization and encryption phase for this variant.

Initialization First, the key, initialization vector (IV) and the constants are loaded into the five registers of the state.

$$\begin{aligned} S_0 &= IV \parallel 0^{128} \\ S_1 &= K \\ S_2 &= 1^{256} \\ S_3 &= 0^{256} \\ S_4 &= \text{constant} \end{aligned}$$

For MORUS-1280-256, K denotes the 256-bit key, for MORUS-1280-128 K is built by the concatenation of the 128-bit key with itself. The *constant* in the 5th register is the Fibonacci sequence modulo 256. After preparing the initial state, the state update function is called 16 times. Each of these 16 steps consists 5 rounds, so there are $5 \times 16 = 80$ rounds in the initialization. Finally, the key is XORed with the second register S_1 .

Encryption Once the initialization is finished, the plaintext block P_i is encrypted.

$$C_i = P_i \oplus S_0 \oplus (S_1 \lll 192) \oplus (S_2 \& S_3)$$

Then, the plaintext P_i is used to update the state. Once the state is updated, the next plaintext block P_{i+1} can be encrypted.

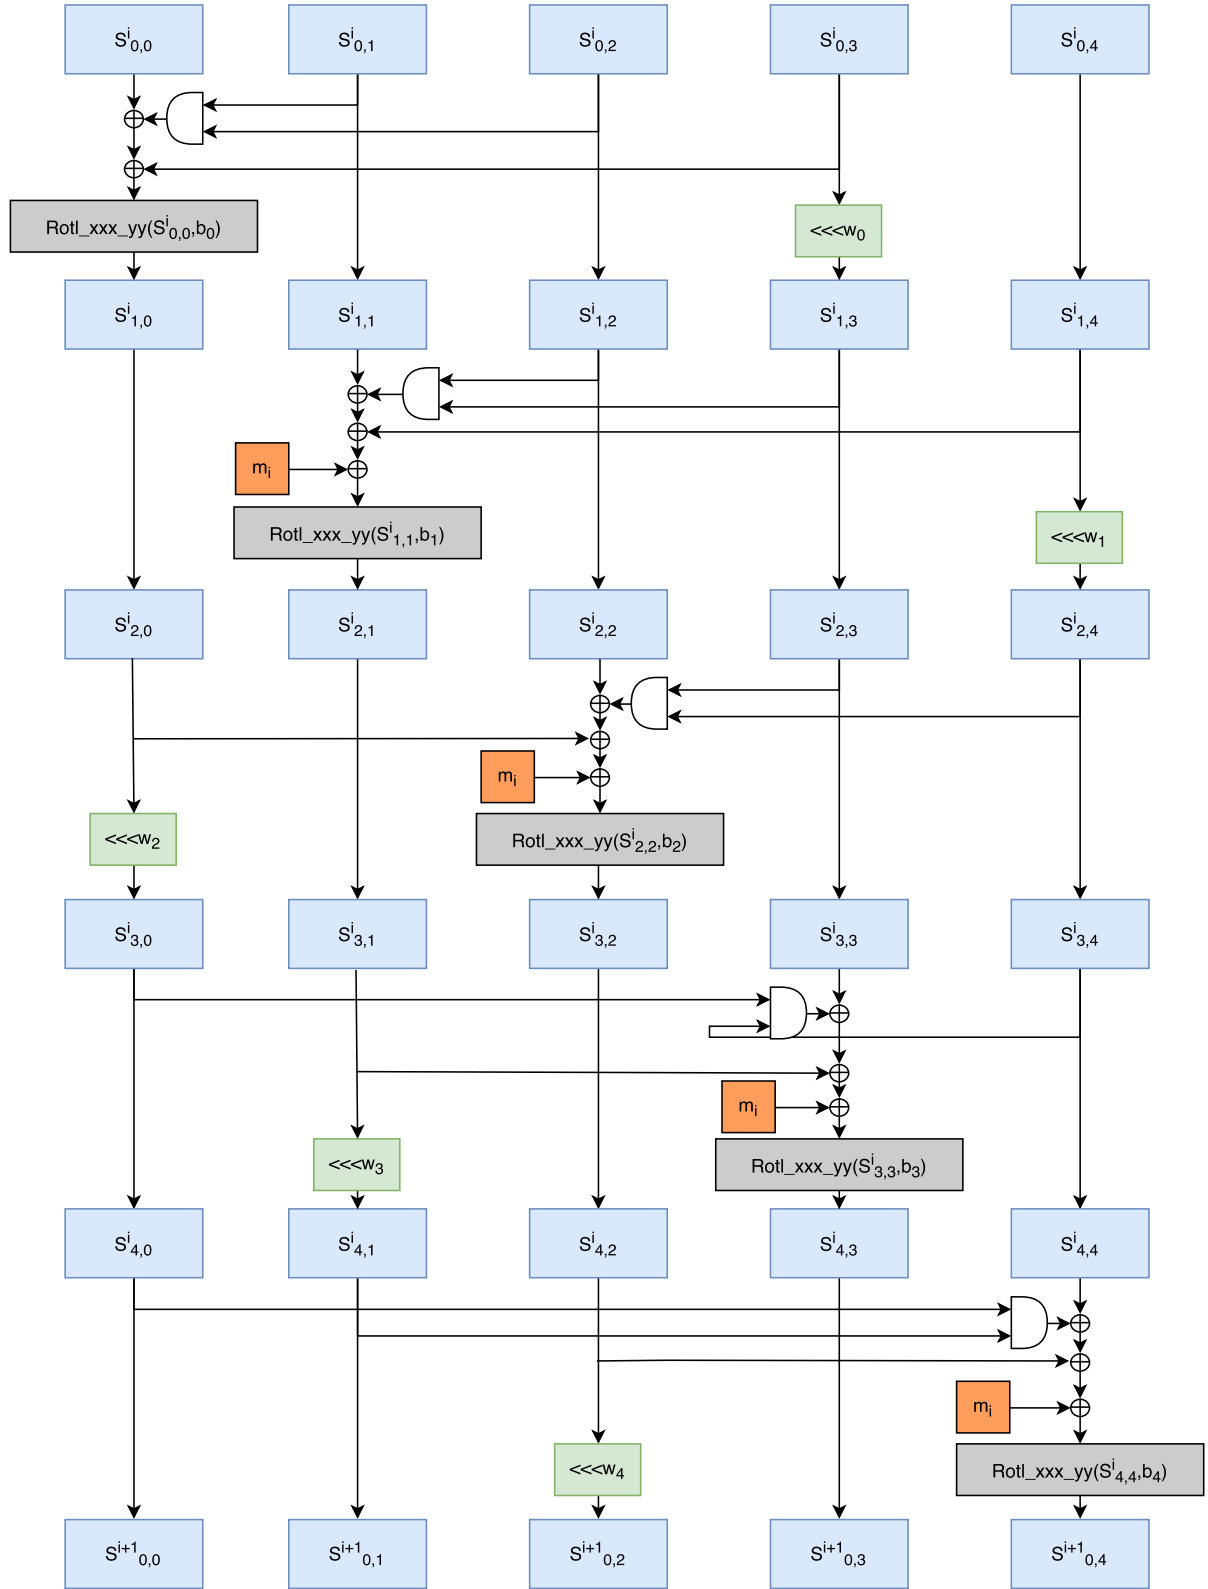


Figure 1: The state update function of MORUS

Table 1: Our attacks on reduced variants of MORUS

Algorithm	Initialization rounds	Technique	Attack type	Complexity	Reference
MORUS-1280-256	18	differential	key recovery	2^{253}	Sect. 3
MORUS-1280-128	7	internal differential	ciphertext prediction	2^1	Sect. 4.1
MORUS-1280-256	8	rotational	key recovery	2^{251}	Sect. 5.1
MORUS-1280-256	10	rotational	ciphertext prediction	2^1	Sect. 5.2

3 KEY RECOVERY ATTACK ON MORUS WITH 18-ROUND INITIALIZATION

A common approach for attacking block ciphers with differential cryptanalysis is the so called $R - 1$ attack. First, the attacker constructs a differential characteristic which covers all but last round. Then, by guessing a round key from the last round, she inverts the last round from given ciphertexts. If the statistics of ‘new’ ciphertexts behaves accordingly to the differential characteristic, then the round key guess is correct. However, MORUS has a completely different structure than a typical modern block cipher (such as AES) and it is not clear how the classic key recovery attacks could be used against MORUS. Also, applying block-cipher-like differential cryptanalysis directly to the encryption phase is not possible because of the nonce requirement.

To circumvent these limitations, we propose a new, different way for the key recovery. Our idea can be viewed as an accelerated exhaustive search. A general scheme of the attack is as follows. For a given 256-bit key, there are 256 differential counterparts which differ from a given key on a single bit. We try to guess/identify one of such counterparts and then we argue that a total cost for this is lower than the exhaustive search. Once the differential counterpart is identified and the XOR difference with the key is known, clearly the secret key itself is also revealed.

The attack below is described for MORUS-1280-256, where the initialization phase is reduced to 18 rounds. For a 256-bit key, the probability that we guess one of its 256 counterparts is $2^{-256} \times 256 = 2^{-248}$, so after 2^{248} guesses we should hit the differential counterpart of the key. Now the question is how to check (cheaply) whether we guess the counterpart. Here help comes from differential characteristics, such as the one shown in Table 2. In the table differences are given in the hexadecimal format. When bits with a non-zero input difference are multiplied (bitwise AND), the output difference depends on the actual values of these bits. Since we do not know these actual values (only their differences), such

an output difference can not be determined with certainty. We denote a nibble with the undetermined difference by ‘?’. (To avoid heavy and unreadable notation, the symbol ‘?’ captures both partly undetermined nibbles and nibbles, where all 4 bits are undetermined.)

In a precomputation phase, we construct 256 differential characteristics, each with 1-bit difference initial state. The 18-round characteristics end with a ciphertext (256 bits), where differences for some bits are still known. These characteristics are stored in memory and then used to identify whether we guess the counterpart of the key or not. Algorithm 1 shows the pseudocode of the attack.

Algorithm 1 Key recovery attack on MORUS-1280-256 with 18-round initialization.

```

1:  $P \leftarrow 0x00 \dots 0$     ▷ Set plaintext to all-zero vector
2:  $IV \leftarrow 0x00 \dots 0$     ▷ Set IV to all-zero vector
3: MORUS-1280-256( $IV, key$ )    ▷ Initialization
4:  $C \leftarrow P \oplus S_0 \oplus (S_1 \lll 192) \oplus (S_2 \& S_3)$     ▷
   Calculate the ciphertext  $C$ 

5: for  $i \leftarrow 0$  to  $i < 2^{248}$  do
6:   MORUS-1280-256( $IV, guessedKey$ )    ▷
   Initialization
7:    $C' \leftarrow P \oplus S_0 \oplus (S_1 \lll 192) \oplus (S_2 \& S_3)$     ▷
   Calculate the ciphertext  $C'$ 
8:   for  $n \leftarrow 1$  to 256 do
9:     if  $(C \oplus C'$  matches the characteristic $_n$ )
       then
10:       $(guessedKey \oplus characteristic_n \text{ input})$ 
        is a candidate for the secret key
11:     end if
12:   end for
13: end for

```

After 2^{248} guesses the attacker should hit one of the 256 differential counterparts of the secret key. Hence the main loop is iterated 2^{248} times. In the inner loop, the ciphertexts difference is calculated and if it matches the output from one of the 256 characteristics (generated in the precomputation phase), then the guessed key might be the counterpart of the se-

Table 2: The 18-round differential characteristic for MORUS-1280-256

Initial state of differences				
S_0	0000000000000000	0000000000000000	0000000000000000	0000000000000000
S_1	0000080000000000	0000000000000000	0000000000000000	0000000000000000
S_2	0000000000000000	0000000000000000	0000000000000000	0000000000000000
S_3	0000000000000000	0000000000000000	0000000000000000	0000000000000000
S_4	0000000000000000	0000000000000000	0000000000000000	0000000000000000
↓				
Step 1 (Rounds 1–5)				
S_0	0000000000000000	0100000000000000	0000000000000000	0000000000000000
S_1	0000000000000000	0000000000000000	0000000002000000	0000000000000000
S_2	0000000000000000	0000000040000000	0000000000000000	0000000000000000
S_3	0000000100000000	0000000000000000	0000000000000000	0000000000000000
S_4	0000000400000000	?0000000000000000	00000000?0000000	0000000000000000
ciphertext: 254 known differences				
↓				
Step 2 (Rounds 6–10)				
S_0	0000000000000000	0000200000000000	00000?0000000020	000000?000000000
S_1	0000000000000?80	000000000000?000	0000000000010000	00000?000000?000
S_2	0000000?000000?0	0000000000008000?	00000800000?0010	000000000000??00
S_3	000000?000800000	00?0000000?0000	000?0?00000?0000	0000?08000?0000?
S_4	0000000?0080?00?	000?800000?0100	0000?04000?0000	0000??00000?0?0?
ciphertext: 229 known differences				
↓				
Step 3 (Rounds 11–15)				
S_0	0??0000?0??0000	000??0100??0000	?400000?0?0?000?	??00000?0?0000
S_1	???0000??0000?	??0?000??00?0	??0000?00??02?	?04?000?0?0000?
S_2	????00??00?00	????0000?0?00	0??0?0?10??0??	??????0?00?0??
S_3	??0??????1??	2?????0??0??	??04???8????	??00????00?0??
S_4	???04?????0??	?????1??????	??????8?????	??00????0?0??
ciphertext: 101 known differences				
↓				
Step 4 (Rounds 16–18)				
S_0	????????????	????????2?????	????????????	?8????????????
S_1	????????????	????????????	????????????	????????????
S_2	????????????	????????????	????????????	????????????
S_3	2?????0?????	??04???8????	??00????00?0??	??0??????1??
S_4	???????8????	??00????0?0??	??04??????0??	?????1??????
ciphertext: 2 known differences				

cret key. Please note that the cost of the inner loop is very small. Basically, we need one bitwise XOR between C and C' and then bitwise AND with the 256-bit mask, which tells which bits should be taken into account. For comparison, the 18-round initialization requires 252 XOR operations on registers and 90 bitwise AND between registers. So, the total cost of the inner loop is negligible as it is just a tiny fraction of a single call to the 18-round MORUS-1280-256.

Outputs (ciphertexts) from the 18-round differential characteristics have, on average, 3 bits with known differences. So, for 256-bit ciphertexts the ‘filter’ is not very strong and we expect around $2^{256}/2^3 = 2^{253}$ false alarms to check. Thus, checking false alarms is a dominant factor of the time complexity of the attack.

We could try to use more differential characteristics which start with a 2- or 3-bit difference. This would lead to fewer iterations in the main loop (lower complexity of the attack), however, our experiments show that we could penetrate fewer rounds in that cases. Indeed, this is expected as more differences in the initial state lead to faster propagation of differences and unknown differences start dominating the state earlier.

4 INTERNAL DIFFERENTIAL CRYPTANALYSIS OF MORUS

In internal differential cryptanalysis, the attacker follows differences between parts of the state. However, it is not obvious which parts should be compared. Thus, the first step is to specify how we calculate an internal difference. A hint for us is that a symmetric state (zero internal difference) should be preserved (ideally) by all the operations in the algorithm. If some operations break the chosen symmetry, then we should look for another symmetry. (Though it might be the case that there is not such a symmetry.) For MORUS there are three ‘natural’ choices how to define the symmetry and calculate internal differences. We could try a difference between two halves of a state, between two halves of a register or between two halves of a word. The symmetry between halves of the state is broken by both \lll and $Rotl$ operations, whereas the other two symmetries are preserved by all the operations. In this analysis we focus on a symmetry within the registers and the reason for this choice is the following. MORUS-1280 operates on 256-bit registers and the 128-bit secret key is concatenated with its copy ($key||key$), then loaded in the S_1 register in the initial state. Thus, there is a symmetry in S_1 , regardless of the actual value of the key. The other

choice of symmetry (within words) is discussed at the end of this section.

4.1 Ciphertext Prediction for MORUS-1280-128 with 7-round Initialization

Let us first analyse a propagation of internal differences for MORUS-1280-128. The IV has to be set to all 0’s to have the first register S_0 symmetric. As explained above, the S_1 register is symmetric, regardless of a key value. The subsequent registers S_2 and S_3 are initialized with all 1’s and all 0’s, so they are clearly symmetric. Finally, the fifth register S_4 is filled with a constant, which is the Fibonacci sequence modulo 256.

Clearly, this is not a symmetric constant, so the initial state has some internal differences (the symmetry is partly disturbed). Table 3 shows an evolution of the internal differences up to 8 rounds. If the initialization is reduced to 7 rounds, the attacker is able to predict a ciphertext bit with probability 1. As shown in Table 3, the ciphertext after 7 rounds still has one known internal difference. Specifically, the difference between bits number 29 and 157 is known. The attacker initializes MORUS-1280-128 with the IV set to all 0’s and obtains the 29th bit from a ciphertext block. Knowing a plaintext and the difference between 29th and 157th bit from a ciphertext, she can predict the value of the unseen 157th ciphertext bit. The prediction works for any key and any known plaintext.

We also investigate the internal differentials with probabilistic transitions. However, we could not find any low-cost differential path, which leads to many known ciphertext differences after 8 rounds.

Other symmetry The other symmetry, which is preserved by both \lll and $Rotl$ operations, is the symmetry between halves of a word. Interestingly, in MORUS a rotation number in \lll is always a multiple of a word size (due to implementation efficiency) and it helps us to keep the symmetry as \lll moves the whole words ‘intact’ to a new position in the state. However, using the symmetry within words for the ciphertext prediction is more difficult. To have the S_1 register (initialized with 128-bit key) symmetric, we have to assume that a key itself is symmetric. Therefore the analysis is valid only against the weak key class of 2^{64} symmetric keys. We investigate how the internal differences evolve for this symmetry and the result is very similar to findings from Table 3, that is up to 7 rounds there are known differences in the ciphertext (see Appendix).

Table 3: The 8-round internal differential characteristic for MORUS-1280-128

Initial state of internal differences				
S_0	00000000	00000000	00000000	00000000
S_1	00000000	00000000	00000000	00000000
S_2	00000000	00000000	00000000	00000000
S_3	00000000	00000000	00000000	00000000
S_4	db3c1957	6ec727fc	3533061b	e35c51bf

↓
Round 1

S_0	00000000	00000000	00000000	00000000
S_1	00000000	00000000	00000000	00000000
S_2	00000000	00000000	00000000	00000000
S_3	00000000	00000000	00000000	00000000
S_4	db3c1957	6ec727fc	3533061b	e35c51bf

ciphertext: 128 known internal differences

↓
Round 2

S_0	00000000	00000000	00000000	00000000
S_1	ed9e0cab	376393fe	9a99830d	f1ae28df
S_2	00000000	00000000	00000000	00000000
S_3	00000000	00000000	00000000	00000000
S_4	3533061b	e35c51bf	db3c1957	6ec727fc

ciphertext: 128 known internal differences

↓
Round 3

S_0	00000000	00000000	00000000	00000000
S_1	ed9e0cab	376393fe	9a99830d	f1ae28df
S_2	00000000	00000000	00000000	00000000
S_3	00000000	00000000	00000000	00000000
S_4	3533061b	e35c51bf	db3c1957	6ec727fc

ciphertext: 128 known internal differences

↓
Round 4

S_0	00000000	00000000	00000000	00000000
S_1	9a99830d	f1ae28df	ed9e0cab	376393fe
S_2	00000000	00000000	00000000	00000000
S_3	????????	????????	??????6?	????????
S_4	3533061b	e35c51bf	db3c1957	6ec727fc

ciphertext: 27 known internal differences

↓

Round 5

↓

S_0	00000000	00000000	00000000	00000000
S_1	9a99830d	f1ae28df	ed9e0cab	376393fe
S_2	00000000	00000000	00000000	00000000
S_3	????????	????????	??????6?	????????
S_4	?0??????	????????	????????	????????

ciphertext: 27 known internal differences

↓
Round 6

↓

S_0	????????	????????	????????	????????
S_1	9a99830d	f1ae28df	ed9e0cab	376393fe
S_2	00000000	00000000	00000000	00000000
S_3	????????	??????6?	????????	????????
S_4	?0??????	????????	????????	????????

ciphertext: 4 known internal differences

↓
Round 7

↓

S_0	????????	????????	????????	????????
S_1	????????	????????	????????	????????
S_2	00000000	00000000	00000000	00000000
S_3	????????	??????6?	????????	????????
S_4	????????	????????	?0??????	????????

ciphertext: 1 known internal difference

↓
Round 8

↓

S_0	????????	????????	????????	????????
S_1	????????	????????	????????	????????
S_2	????????	????????	????????	????????
S_3	????????	??????6?	????????	????????
S_4	????????	????????	?0??????	????????

ciphertext: 0 known internal differences

5 ROTATIONAL CRYPTANALYSIS OF MORUS

In this section we show how to use rotational cryptanalysis to mount the key recovery attack and the ciphertext prediction for MORUS with the round-reduced initialization phase. We study how rotational relations between two states change over the subsequent rounds and then try to use these observations for the attacks.

First, we should specify what we mean by rotational relation between two MORUS states. Having two states S and S' , the state S' is rotated (with a reference to S) if all words in S' are copied from S and rotated by a rotational number n . (In MORUS-1280, a word size is 64 bits, so there are 64 possible values of n .) With S' constructed like this, all the steps in the algorithm preserve the rotational relation between S and S' . Similarly to internal differentials, where the constant breaks the symmetry, here the constants breaks the desired rotational relation. The pseudorandom constant (Fibonacci sequence) breaks the rotational pattern regardless of a chosen value of n .

Now, let us see an example of the rotational characteristic for MORUS-1280-256. We use the same notation as for internal differentials but this time a hexadecimal number means the difference between corresponding bits from the states S and S' . So, for example, if we consider the 1st bit and rotational number $n = 8$, then we calculate the XOR difference between 1st bit from S and 9th bit from S' (within a given word). Table 4 shows the rotational characteristic¹ with $n = 43$. The registers $S_0 \dots S_3$ have a (rotational) differences set to 0, whereas the 5th register is initialized with the Fibonacci constant, hence there are differences and rotational relation is disturbed.

5.1 Key-recovery Attack on MORUS with 8-round Initialization

In the following key recovery attack we use our idea of accelerated exhaustive search from Section 3. This time, however, we take advantage of rotational characteristics rather than differential. The attack is applied to MORUS-1280-256 with the initialization phase reduced to 8 rounds. The general scheme of the attack is similar to the differential one, yet some details specific rotational analysis need to be explained.

The main idea behind the attack is that we can find the rotational counterpart of the secret key with

¹Please note that here we trace differences between two states, so we need two times more symbols than in the case of internal differences of a single state.

an effort lower than the exhaustive search. Once the rotational counterpart is found, we rotate it back to get the key. The attack description given below is for MORUS-1280-256 but the attack works in the same way for MORUS-1280-128.

A register in MORUS-1280 has 64-bit words. Therefore, there are 64 possible values of the rotational number n , including the identity transformation. For 256-bit key, the probability that we guess one of its rotational counterparts is $2^{-256} \times 64 = 2^{-250}$, so after 2^{250} guesses we should hit the rotational counterpart of the key². Now the question is how to check (cheaply) whether we guess the rotational counterpart. Here the rotational characteristics (such as the one shown in Table 4) come in handy. In Precomputation, we generate 64 rotational characteristics, each corresponds to a different rotational number. The 8-round characteristics end with a ciphertext (256 bits), where rotational relations for some bits are still known. These characteristics are stored in memory and then used to identify whether we guess the rotational counterpart of the key or not. Algorithm 2 shows the pseudocode of the attack.

Algorithm 2 Key recovery attack on MORUS-1280-256 with 8-round initialization.

```

1:  $P \leftarrow 0x00 \dots 0$     ▷ Set plaintext to all-zero vector
2:  $IV \leftarrow 0x00 \dots 0$     ▷ Set IV to all-zero vector
3: MORUS-1280-256( $IV, key$ )    ▷ Initialization
4:  $C \leftarrow P \oplus S_0 \oplus (S_1 \lll 192) \oplus (S_2 \& S_3)$     ▷
   Calculate the ciphertext  $C$ 

5: for  $i \leftarrow 0$  to  $i < 2^{250}$  do
6:   MORUS-1280-256( $IV, guessedKey$ )    ▷
   Initialization
7:    $C' \leftarrow P \oplus S_0 \oplus (S_1 \lll 192) \oplus (S_2 \& S_3)$     ▷
   Calculate the ciphertext  $C'$ 
8:   for  $n \leftarrow 0$  to 64 do
9:     if  $(C \oplus C'$  matches the characteristic $_n)$ 
       then
10:       $(guessedKey \lll n)$  is a candidate for
        the secret key
11:     end if
12:   end for
13: end for

```

As explained above, after 2^{250} the attacker should

²There is a class of symmetric keys, where both halves of a word is the same. For such keys, rotations act as the identity function hence there are not any rotational counterparts. However, 2^{128} such keys is only a tiny fraction of all 2^{256} possible keys and checking the whole class before the actual attack does not affect the total complexity of our key recovery attack.

Table 4: The 9-round rotational characteristic ($n = 43$) for MORUS-1280-256

Initial state		Round 5	
		↓	
S_0	0000000000000000 0000000000000000	S_0	0000000000000000 0000000000000000
S_1	0000000000000000 0000000000000000	S_1	2127f290a5a32140 ced02255f610be8f
S_2	0000000000000000 0000000000000000	S_2	46054a105a4082c3 ed7957ba49fc2054
S_3	0000000000000000 0000000000000000	S_3	02a5082d204161a3 bcabdd24fe102a76
S_4	284169020b0d1815 5ee927f08153b5e5	S_4	93f94852d190a010 68112afb085f47e7
	ca42968c8500849f 8957d842fa3f3b40		a42968c8500849fc 957d842fa3f3b408
		ciphertext: 146 known differences	
↓		Round 6	
Round 1		↓	
S_0	0000000000000000 0000000000000000	S_0	??????????6??? ??????????????
S_1	0000000000000000 0000000000000000	S_1	2127f290a5a32140 ced02255f610be8f
S_2	0000000000000000 0000000000000000	S_2	46054a105a4082c3 ed7957ba49fc2054
S_3	0000000000000000 0000000000000000	S_3	02a5082d204161a3 bcabdd24fe102a76
S_4	284169020b0d1815 5ee927f08153b5e5	S_4	93f94852d190a010 68112afb085f47e7
	ca42968c8500849f 8957d842fa3f3b40		a42968c8500849fc 957d842fa3f3b408
ciphertext: 256 known differences		ciphertext: 90 known differences	
↓		Round 7	
Round 2		↓	
S_0	0000000000000000 0000000000000000	S_0	??????????6??? ??????????????
S_1	46054a105a4082c3 ed7957ba49fc2054	S_1	2127f290a5a32140 ced02255f610be8f
S_2	2127f290a5a32140 ced02255f610be8f	S_2	46054a105a4082c3 ed7957ba49fc2054
S_3	0000000000000000 0000000000000000	S_3	02a5082d204161a3 bcabdd24fe102a76
S_4	0000000000000000 0000000000000000	S_4	93f94852d190a010 68112afb085f47e7
	ca42968c8500849f 8957d842fa3f3b40		a42968c8500849fc 957d842fa3f3b408
ciphertext: 256 known differences		ciphertext: 48 known differences	
↓		Round 8	
Round 3		↓	
S_0	0000000000000000 0000000000000000	S_0	??????????f???? ???????????6???
S_1	46054a105a4082c3 ed7957ba49fc2054	S_1	2127f290a5a32140 ced02255f610be8f
S_2	2127f290a5a32140 ced02255f610be8f	S_2	46054a105a4082c3 ed7957ba49fc2054
S_3	0000000000000000 0000000000000000	S_3	02a5082d204161a3 bcabdd24fe102a76
S_4	0000000000000000 0000000000000000	S_4	93f94852d190a010 68112afb085f47e7
	ca42968c8500849f 8957d842fa3f3b40		a42968c8500849fc 957d842fa3f3b408
ciphertext: 256 known differences		ciphertext: 20 known differences	
↓		Round 9	
Round 4		↓	
S_0	0000000000000000 0000000000000000	S_0	??????????f???? ???????????6???
S_1	46054a105a4082c3 ed7957ba49fc2054	S_1	2127f290a5a32140 ced02255f610be8f
S_2	2127f290a5a32140 ced02255f610be8f	S_2	46054a105a4082c3 ed7957ba49fc2054
S_3	0000000000000000 0000000000000000	S_3	02a5082d204161a3 bcabdd24fe102a76
S_4	0000000000000000 0000000000000000	S_4	93f94852d190a010 68112afb085f47e7
	ca42968c8500849f 8957d842fa3f3b40		a42968c8500849fc 957d842fa3f3b408
ciphertext: 256 known differences		ciphertext: 1 known difference	
↓			

guess one of the 64 rotational counterparts of the secret key. Hence the main loop is iterated 2^{250} times. In the inner loop, the ciphertexts difference is calculated and if it matches one of the 64 characteristics (generated in Precomputation), then the guessed key (rotated by n) might be the secret key. Please note that the cost of the inner loop is very small. Basically, we need one bitwise XOR between C and C' and then bitwise AND with the 256-bit mask, which tells which bits should be taken into account. For comparison, the 8-round initialization requires 112 XOR operations on registers and 40 bitwise AND between registers. So, the total cost of the inner loop is not more than a single call to the 8-round MORUS-1280.

Outputs (ciphertexts) from the 8-round rotational characteristics have, on average, 11 bits with known differences. So, for 256-bit ciphertexts the ‘filter’ is not very strong and there will be many false alarms. For example, when the rotational number n is 43, there are 20 known differences in the ciphertext, hence we expect around $2^{250}/2^{20} = 2^{230}$ false positives. We calculate an expected number of false alarms for all 64 values of n and the total number of false alarms (candidates for one of the rotational counterpart of the secret key) is about 2^{248} . Therefore, the time complexity of the attack is dominated by calls to MORUS in the main loop (2^{250} iterations). Adding false alarms (2^{248}) plus a small cost of bitwise operations in the inner loop make the total cost of the attack equal to around 2^{251} .

More rounds There are rotational characteristics (such as the one shown in Table 4), which cover 9 or even 10 rounds. However, a number of known differences in the ciphertexts (for 9 or 10 rounds) is, on average, less than 1. This means that a number of false alarms would be very close to 2^{256} and it would be difficult to argue that the total cost of the attack is lower than the exhaustive search for the 256-bit key.

5.2 Ciphertext Prediction in Weak Key Scenario for 10-round MORUS

When we assume that the 256-bit secret key is one of 2^{128} symmetric keys (for which halves of a given word are the same), then we could use the rotational characteristics to predict the ciphertext, similarly to the attack described in Section 4.1. For the 10-round rotational characteristic, where $n = 18$, there are 4 known differences between the ciphertext and its rotational counterpart. The attacker initializes MORUS-1280-256 with an arbitrarily chosen IV and obtains the ciphertext block. Now the attacker can predict the ciphertext (4 bits), which would be generated from

MORUS initialized with the rotated IV (each word of the IV rotated by 18). So, using rotational characteristics the prediction is possible up to 10 rounds, yet we need a stronger assumption — a key from the weak key class.

6 Conclusion

We have proposed a new approach for the theoretical key recovery attack against the round-reduced MORUS. The technique can be seen as an accelerated exhaustive search and it works not only with differential cryptanalysis but also with other types of distinguishers. The technique could be particularly useful for ciphers, which have completely different structure than typical block ciphers such as AES, for example for the sponge-based cryptographic primitives (Bertoni et al., a).

We have also analysed the resistance of MORUS against internal differentials and rotational cryptanalysis. Our findings have revealed that the cipher offers solid security margin against these techniques. As MORUS has some unorthodox design features (such as a lack of round constants) we think it is essential to analyse such new, promising algorithms with a possibly wide range of cryptanalytic tools and techniques. All performed test are applicable for old and new version of submitted cipher. Our work helps to realize this goal.

REFERENCES

- CAESAR: Competition for Authenticated Encryption: Security, Applicability, and Robustness. <http://competitions.cr.yp.to/caesar.html>.
- Bernstein, D. J. (2005). Salsa20. Technical report, eSTREAM, ECRYPT Stream Cipher Project. <http://cr.yp.to/snuffle.html>.
- Bertoni, G., Daemen, J., Peeters, M., and Van Assche, G. Cryptographic Sponges. <http://sponge.noekeon.org/CSF-0.1.pdf>.
- Bertoni, G., Daemen, J., Peeters, M., and Van Assche, G. Keccak Sponge Function Family Main Document. <http://keccak.noekeon.org/Keccak-main-2.1.pdf>.
- Dinur, I., Dunkelman, O., and Shamir, A. (2013). Collision Attacks on Up to 5 Rounds of SHA-3 Using Generalized Internal Differentials. In *Fast Software Encryption - 20th International Workshop, FSE 2013, Singapore, March 11-13, 2013. Revised Selected Papers*, pages 219–240.
- Dwivedi, A. D., Kloucek, M., Morawiecki, P., Nikolic, I., Pieprzyk, J., and Wójtowicz, S. (2016). Sat-based cryptanalysis of authenticated ciphers from the CAESAR competition. *IACR Cryptology ePrint Archive*, 2016:1053.
- Jean, J. and Nikolic, I. (2015). Internal differential boomerangs: Practical analysis of the round-reduced keccak-f permutation. In Leander, G., editor, *FSE 2015*, volume 9054 of *LNCS*, pages 537–556. Springer, Heidelberg.
- Khovratovich, D. and Nikolić, I. (2010). Rotational cryptanalysis of ARX. In *Proceedings of the 17th international conference on Fast software encryption*, *LNCS*, pages 333–346. Springer-Verlag.
- Khovratovich, D., Nikolic, I., and Rechberger, C. (2010). Rotational Rebound Attacks on Reduced Skein. In *ASIACRYPT'10*, volume 6477 of *LNCS*, pages 1–19.
- Knudsen, L. R., Matusiewicz, K., and Thomsen, S. S. (2009). Observations on the Shabal keyed permutation. Available online. <http://www.mat.dtu.dk/people/S.Thomsen/shabal/shabal.pdf>.
- Mileva, A., Dimitrova, V., and Velichkov, V. (2015). Analysis of the authenticated cipher MORUS (v1). In *Cryptography and Information Security in the Balkans - Second International Conference, BalkanCryptSec 2015, Koper, Slovenia, September 3-4, 2015, Revised Selected Papers*, pages 45–59.
- Morawiecki, P., Pieprzyk, J., and Srebrny, M. (2013). Rotational cryptanalysis of round-reduced Keccak. In *Fast Software Encryption*, *LNCS*. Springer.
- MORUSGoogleGroupDiscussin. <https://groups.google.com/forum/#!msg/crypto-competitions/p1TQVraGkrU/C0bpcrLxMQAJ/>.
- Peyrin, T. (2010). Improved Differential Attacks for ECHO and Grøstl. In *Advances in Cryptology - CRYPTO 2010, 30th Annual Cryptology Conference, Santa Barbara, CA, USA, August 15-19, 2010. Proceedings*, pages 370–392.
- Saarinen, M. O. (2016). The BRUTUS automatic cryptanalytic framework - testing CAESAR authenticated encryption candidates for weaknesses. *J. Cryptographic Engineering*, 6(1):75–82.
- Standaert, F.-X., Piret, G., Gershenfeld, N., and Quisquater, J.-J. (2006). SEA: A Scalable Encryption Algorithm for Small Embedded Applications. In *CARDIS'06*, volume 3928 of *LNCS*, pages 222–236.
- Wu, H. and Huang, T. The Authenticated Cipher MORUS. <https://competitions.cr.yp.to/caesar-submissions.html>.